

Analysis of Yelp Business Intelligence Data

We will analyze a subset of Yelp's business, reviews and user data. This dataset comes to us from [Kaggle](#) although we have taken steps to pull this data into a public s3 bucket: `s3://sta9760-yelpdataset/yelp-light/*business.json`

Installation and Initial Setup

Begin by installing the necessary libraries that you may need to conduct your analysis. At the very least, you must install `pandas` and `matplotlib`

In [1]:

```
#try:
#    print('Installing pip:')
#    sc.install_pypi_package("pip==21.3.1")
#except Exception as e:
#    print(e)

#try:
#    print('Installing numpy:')
#    sc.install_pypi_package("numpy==1.18")
#except Exception as e:
#    print(e)

#try:
#    print('Installing pyspark:')
#    sc.install_pypi_package("pyspark==2.0.0")
#except Exception as e:
#    print(e)

try:
    print('Installing matplotlib:')
    sc.install_pypi_package("matplotlib==3.2.1")
except Exception as e:
    print(e)

try:
    print('Installing pyparsing:')
    sc.install_pypi_package("pyparsing==2.4.7")
except Exception as e:
    print(e)
```

```

try:
    print('Installing python-dateutil:')
    sc.install_pypi_package("python-dateutil==2.8.1")
except Exception as e:
    print('python-dateutil:')
    print(e)

try:
    print('Installing kiwisolver:')
    sc.install_pypi_package("kiwisolver==1.2.0")
except Exception as e:
    print(e)

try:
    print('Installing cyclor:')
    sc.install_pypi_package("cyclor==0.10.0")
except Exception as e:
    print(e)

try:
    print('Installing pandas:')
    sc.install_pypi_package("pandas==1.0.3")
except Exception as e:
    print(e)

try:
    print('Installing pybind11:')
    sc.install_pypi_package("pybind11==2.8.1")
except Exception as e:
    print(e)

try:
    print('Installing scipy:')
    sc.install_pypi_package("scipy==1.4.1")
except Exception as e:
    print(e)

try:
    print('Installing seaborn:')
    sc.install_pypi_package("seaborn==0.11.2")
except Exception as e:
    print(e)

```

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
0	application_1638245194643_0001	pyspark	idle	Link	Link	✓

SparkSession available as 'spark'.

Installing matplotlib:

Collecting matplotlib==3.2.1

Downloading https://files.pythonhosted.org/packages/b2/c2/71fcf957710f3ba1f09088b35776a799ba7dd95f7c2b195ec800933b276b/matplotlib-3.2.1-cp37-cp37m-manylinux1_x86_64.whl (12.4MB)

Collecting python-dateutil>=2.1 (from matplotlib==3.2.1)

Downloading https://files.pythonhosted.org/packages/36/7a/87837f39d0296e723bb9b62bbb257d0355c7f6128853c78955f57342a56d/python_dateutil-2.8.2-py2.py3-none-any.whl (247kB)

Collecting pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 (from matplotlib==3.2.1)

Downloading <https://files.pythonhosted.org/packages/a0/34/895006117f6fce0b4de045c87e154ee4a20c68ec0a4c9a36d900888fb6bc/pyparsing-3.0.6-py3-none-any.whl> (97kB)

Collecting cycler>=0.10 (from matplotlib==3.2.1)

Downloading <https://files.pythonhosted.org/packages/5c/f9/695d6bedebd747e5eb0fe8fad57b72fdf25411273a39791cde838d5a8f51/cycler-0.11.0-py3-none-any.whl>

Requirement already satisfied: numpy>=1.11 in /usr/local/lib64/python3.7/site-packages (from matplotlib==3.2.1)

Collecting kiwisolver>=1.0.1 (from matplotlib==3.2.1)

Downloading https://files.pythonhosted.org/packages/09/6b/6e567cb2e86d4e5939a9233f8734e26021b6a9c1bc4b1edccba236a84cc2/kiwisolver-1.3.2-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.whl (1.1MB)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.1->matplotlib==3.2.1)

Installing collected packages: python-dateutil, pyparsing, cycler, kiwisolver, matplotlib

Successfully installed cycler-0.11.0 kiwisolver-1.3.2 matplotlib-3.2.1 pyparsing-3.0.6 python-dateutil-2.8.2

Installing pyparsing:

Collecting pyparsing==2.4.7

Downloading <https://files.pythonhosted.org/packages/8a/bb/488841f56197b13700afd5658fc279a2025a39e22449b7cf29864669b15d/pyparsing-2.4.7-py2.py3-none-any.whl> (67kB)

Installing collected packages: pyparsing

Found existing installation: pyparsing 3.0.6

Uninstalling pyparsing-3.0.6:

Successfully uninstalled pyparsing-3.0.6

Successfully installed pyparsing-2.4.7

Installing python-dateutil:

Collecting python-dateutil==2.8.1

Downloading https://files.pythonhosted.org/packages/d4/70/d60450c3dd48ef87586924207ae8907090de0b306af2bce5d134d78615cb/python_dateutil-2.8.1-py2.py3-none-any.whl (227kB)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil==2.8.1)

Installing collected packages: python-dateutil

Found existing installation: python-dateutil 2.8.2

Uninstalling python-dateutil-2.8.2:
Successfully uninstalled python-dateutil-2.8.2
Successfully installed python-dateutil-2.8.1

Installing kiwisolver:
Collecting kiwisolver==1.2.0
Downloading https://files.pythonhosted.org/packages/31/b9/6202dcae729998a0ade30e80ac00f616542ef445b088ec970d407dfd41c0/kiwisolver-1.2.0-cp37-cp37m-manylinux1_x86_64.whl (88kB)
Installing collected packages: kiwisolver
Found existing installation: kiwisolver 1.3.2
Uninstalling kiwisolver-1.3.2:
Successfully uninstalled kiwisolver-1.3.2
Successfully installed kiwisolver-1.2.0

Installing cyclcr:
Collecting cyclcr==0.10.0
Downloading <https://files.pythonhosted.org/packages/f7/d2/e07d3ebb2bd7af696440ce7e754c59dd546ffe1bbe732c8ab68b9c834e61/cyclcr-0.10.0-py2.py3-none-any.whl>
Requirement already satisfied: six in /usr/local/lib/python3.7/site-packages (from cyclcr==0.10.0)
Installing collected packages: cyclcr
Found existing installation: cyclcr 0.11.0
Uninstalling cyclcr-0.11.0:
Successfully uninstalled cyclcr-0.11.0
Successfully installed cyclcr-0.10.0

Installing pandas:
Collecting pandas==1.0.3
Downloading https://files.pythonhosted.org/packages/4a/6a/94b219b8ea0f2d580169e85ed1edc0163743f55aaeca8a44c2e8fc1e344e/pandas-1.0.3-cp37-cp37m-manylinux1_x86_64.whl (10.0MB)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas==1.0.3)
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from pandas==1.0.3)
Requirement already satisfied: python-dateutil>=2.6.1 in /mnt/tmp/1638246474516-0/lib/python3.7/site-packages (from pandas==1.0.3)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas==1.0.3)
Installing collected packages: pandas
Successfully installed pandas-1.0.3

Installing pybind11:
Collecting pybind11==2.8.1
Downloading <https://files.pythonhosted.org/packages/a8/3b/fc246e1d4c7547a7a07df830128e93c6215e9b93dcb118b2a47a70726153/pybind11-2.8.1-py2.py3-none-any.whl> (208kB)
Installing collected packages: pybind11
Successfully installed pybind11-2.8.1

```

Installing scipy:
Collecting scipy==1.4.1
  Downloading https://files.pythonhosted.org/packages/dd/82/c1fe128f3526b128cfd185580ba40d01371c5d299fcf7f77968e22dfcc2e/scipy-1.4.1-cp37-cp37m-manylinux1_x86_64.whl (26.1MB)
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from scipy==1.4.1)
Installing collected packages: scipy
Successfully installed scipy-1.4.1

Installing seaborn:
Collecting seaborn==0.11.2
  Downloading https://files.pythonhosted.org/packages/10/5b/0479d7d845b5ba410ca702ffcd7f2cd95a14a4dfff1fde2637802b258b9b/seaborn-0.11.2-py3-none-any.whl (292kB)
Requirement already satisfied: numpy>=1.15 in /usr/local/lib64/python3.7/site-packages (from seaborn==0.11.2)
Requirement already satisfied: scipy>=1.0 in /mnt/tmp/1638246474516-0/lib/python3.7/site-packages (from seaborn==0.11.2)
Requirement already satisfied: matplotlib>=2.2 in /mnt/tmp/1638246474516-0/lib/python3.7/site-packages (from seaborn==0.11.2)
Requirement already satisfied: pandas>=0.23 in /mnt/tmp/1638246474516-0/lib/python3.7/site-packages (from seaborn==0.11.2)
Requirement already satisfied: python-dateutil>=2.1 in /mnt/tmp/1638246474516-0/lib/python3.7/site-packages (from matplotlib>=2.2->seaborn==0.11.2)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /mnt/tmp/1638246474516-0/lib/python3.7/site-packages (from matplotlib>=2.2->seaborn==0.11.2)
Requirement already satisfied: cyclor>=0.10 in /mnt/tmp/1638246474516-0/lib/python3.7/site-packages (from matplotlib>=2.2->seaborn==0.11.2)
Requirement already satisfied: kiwisolver>=1.0.1 in /mnt/tmp/1638246474516-0/lib/python3.7/site-packages (from matplotlib>=2.2->seaborn==0.11.2)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas>=0.23->seaborn==0.11.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.1->matplotlib>=2.2->seaborn==0.11.2)
Installing collected packages: seaborn
Successfully installed seaborn-0.11.2

```

```

In [2]: # Confirm our packages have been downloaded
        sc.list_packages()

```

Package	Version
-----	-----
beautifulsoup4	4.9.1
boto	2.49.0
click	7.1.2
cycler	0.10.0
jmespath	0.10.0

joblib	0.16.0
kiwisolver	1.2.0
lxml	4.5.2
matplotlib	3.2.1
mysqlclient	1.4.2
nltk	3.5
nose	1.3.4
numpy	1.16.5
pandas	1.0.3
pip	9.0.1
py-dateutil	2.2
pybind11	2.8.1
pyparsing	2.4.7
python-dateutil	2.8.1
python37-sagemaker-pyspark	1.4.0
pytz	2020.1
PyYAML	5.3.1
regex	2020.7.14
scipy	1.4.1
seaborn	0.11.2
setuptools	28.8.0
six	1.13.0
soupsieve	1.9.5
tqdm	4.48.2
wheel	0.29.0
windmill	1.6

Importing

Now, import the installed packages from the previous block below.

In [3]:

```
# Import packages

import pandas as pd
import pyspark.sql.functions as psf
import seaborn as sns
import matplotlib.pyplot as plt
```

Loading Data

We are finally ready to load data. Using `spark` load the data from S3 into a `dataframe` object that we can manipulate further down in our analysis.

```
In [4]: # Load business data  
df_business = spark.read.json('s3://yelpdatasta9760/yelp_academic_dataset_business.json')
```

Overview of Data

Display the number of rows and columns in our dataset.

```
In [5]: print(f'Total Columns: {len(df_business.dtypes)}')  
print(f'Total Rows: {df_business.count():,}')
```

```
Total Columns: 14  
Total Rows: 160,585
```

Display the DataFrame schema below.

```
In [6]: df_business.printSchema()
```

```
root  
|-- address: string (nullable = true)  
|-- attributes: struct (nullable = true)  
|   |-- AcceptsInsurance: string (nullable = true)  
|   |-- AgesAllowed: string (nullable = true)  
|   |-- Alcohol: string (nullable = true)  
|   |-- Ambience: string (nullable = true)  
|   |-- BYOB: string (nullable = true)  
|   |-- BYOBCorkage: string (nullable = true)  
|   |-- BestNights: string (nullable = true)  
|   |-- BikeParking: string (nullable = true)  
|   |-- BusinessAcceptsBitcoin: string (nullable = true)  
|   |-- BusinessAcceptsCreditCards: string (nullable = true)  
|   |-- BusinessParking: string (nullable = true)  
|   |-- ByAppointmentOnly: string (nullable = true)  
|   |-- Caters: string (nullable = true)
```

```
|  |-- CoatCheck: string (nullable = true)
|  |-- Corkage: string (nullable = true)
|  |-- DietaryRestrictions: string (nullable = true)
|  |-- DogsAllowed: string (nullable = true)
|  |-- DriveThru: string (nullable = true)
|  |-- GoodForDancing: string (nullable = true)
|  |-- GoodForKids: string (nullable = true)
|  |-- GoodForMeal: string (nullable = true)
|  |-- HairSpecializesIn: string (nullable = true)
|  |-- HappyHour: string (nullable = true)
|  |-- HasTV: string (nullable = true)
|  |-- Music: string (nullable = true)
|  |-- NoiseLevel: string (nullable = true)
|  |-- Open24Hours: string (nullable = true)
|  |-- OutdoorSeating: string (nullable = true)
|  |-- RestaurantsAttire: string (nullable = true)
|  |-- RestaurantsCounterService: string (nullable = true)
|  |-- RestaurantsDelivery: string (nullable = true)
|  |-- RestaurantsGoodForGroups: string (nullable = true)
|  |-- RestaurantsPriceRange2: string (nullable = true)
|  |-- RestaurantsReservations: string (nullable = true)
|  |-- RestaurantsTableService: string (nullable = true)
|  |-- RestaurantsTakeOut: string (nullable = true)
|  |-- Smoking: string (nullable = true)
|  |-- WheelchairAccessible: string (nullable = true)
|  |-- WiFi: string (nullable = true)
|-- business_id: string (nullable = true)
|-- categories: string (nullable = true)
|-- city: string (nullable = true)
|-- hours: struct (nullable = true)
|  |-- Friday: string (nullable = true)
|  |-- Monday: string (nullable = true)
|  |-- Saturday: string (nullable = true)
|  |-- Sunday: string (nullable = true)
|  |-- Thursday: string (nullable = true)
|  |-- Tuesday: string (nullable = true)
|  |-- Wednesday: string (nullable = true)
|-- is_open: long (nullable = true)
|-- latitude: double (nullable = true)
|-- longitude: double (nullable = true)
|-- name: string (nullable = true)
|-- postal_code: string (nullable = true)
|-- review_count: long (nullable = true)
|-- stars: double (nullable = true)
|-- state: string (nullable = true)
```


Display the first 5 rows with the following columns:

- business_id
- name
- city
- state
- categories

In [7]:

```
# Show only 6 columns
df_preview = df_business.select('business_id', 'name', 'city', 'state', 'stars', 'categories')
df_preview.show(5)
```

business_id	name	city	state	stars	categories
6iYb2HFDywm3zjuRg...	Oskar Blues Taproom	Boulder	CO	4.0	Gastropubs, Food,...
tCbdrRPZA0oiIYSmH...	Flying Elephants ...	Portland	OR	4.0	Salad, Soup, Sand...
bvN78f1M8NLprQ1a1...	The Reclaimory	Portland	OR	4.5	Antiques, Fashion...
oaepsyvc0J17qwi8c...	Great Clips	Orange City	FL	3.0	Beauty & Spas, Ha...
PE9uqAjdW0E4-8mjG...	Crossfit Terminus	Atlanta	GA	4.0	Gyms, Active Life...

only showing top 5 rows

Analyzing Categories

Let's now answer this question: **how many unique categories are represented in this dataset?**

Essentially, we have the categories per business as a list - this is useful to quickly see what each business might be represented as but it is difficult to easily answer questions such as:

- How many businesses are categorized as `Active Life`, for instance
- What are the top 20 most popular categories available?

Association Table

We need to "break out" these categories from the business ids? One common approach to take is to build an association table mapping a single business id multiple times to each distinct category.

For instance, given the following:

business_id	categories
abcd123	a,b,c

We would like to derive something like:

business_id	category
abcd123	a
abcd123	b
abcd123	c

What this does is allow us to then perform a myriad of rollups and other analysis on this association table which can aid us in answering the questions asked above.

Implement the code necessary to derive the table described from your original yelp dataframe.

```
In [8]: # Build association table for easier analytics of business table
df_t = df_business.withColumn('categories', psf.explode(psf.split("categories", ", ")))
```

```
In [9]: # Select columns for viewing
df_assc = df_t.select('business_id', 'categories')
```

Display the first 5 rows of your association table below.

```
In [10]: df_assc.show(5)
```

```
+-----+-----+
```

```
|      business_id| categories|
+-----+-----+
|6iYb2HFDywm3zjuRg...| Gastropubs|
|6iYb2HFDywm3zjuRg...|      Food|
|6iYb2HFDywm3zjuRg...|Beer Gardens|
|6iYb2HFDywm3zjuRg...| Restaurants|
|6iYb2HFDywm3zjuRg...|      Bars|
+-----+-----+
only showing top 5 rows
```

Total Unique Categories

Finally, we are ready to answer the question: **what is the total number of unique categories available?**

Below, implement the code necessary to calculate this figure.

```
In [11]: # Count of unique categories
df_assc.select(psf.countDistinct("categories")).collect()
```

```
[Row(count(DISTINCT categories)=1330)]
```

Top Categories By Business

Now let's find the top categories in this dataset by rolling up categories.

So now, let's unroll our distinct count a bit and display the per count value of businesses per category.

The expected output should be:

category	count
a	15
b	2
c	45

Or something to that effect.

```
In [12]:
```

```
# Count of businesses in each category
groupby_count = df_assc.groupBy("categories").count().orderBy('count', ascending = False)
groupby_count.show()
```

```
+-----+-----+
|      categories|count|
+-----+-----+
|      Restaurants|50763|
|           Food|29469|
|       Shopping|26205|
|    Beauty & Spas|16574|
|    Home Services|16465|
|Health & Medical|15102|
|    Local Services|12192|
|       Nightlife|11990|
|           Bars|10741|
|    Automotive|10119|
|Event Planning & ...| 9644|
|    Active Life| 9231|
|    Coffee & Tea| 7725|
|    Sandwiches| 7272|
|         Fashion| 6599|
|American (Traditi...| 6541|
|    Hair Salons| 5900|
|         Pizza| 5756|
|    Hotels & Travel| 5703|
| Breakfast & Brunch| 5505|
+-----+-----+
only showing top 20 rows
```

Bar Chart of Top Categories

With this data available, let us now build a barchart of the top 20 categories.

HINT: don't forget about the matplotlib magic!

```
%matplotlib plt
```

In [13]:

```
import matplotlib

# Create barchart vizualization
```


```

pandasDF = groupby_count.toPandas()
pandasDFcharting = pandasDF.head(20)

# Plot
plt.figure(figsize = (15,8))
ax=sns.barplot(x='categories', y='count', data=pandasDFcharting)
ax.set_xticklabels(ax.get_xticklabels(),rotation = 30)

%matplotlib plt

```



Do Yelp Reviews Skew Negative?

Oftentimes, it is said that the only people who write a written review are those who are extremely *dissatisfied* or extremely *satisfied* with the service received.

How true is this really? Let's try and answer this question.

Loading User Data

Begin by loading the user data set from S3 and printing schema to determine what data is available.

```
In [14]: df_reviews = spark.read.json('s3://yelpdatasta9760/yelp_academic_dataset_review.json')

print(f'Total Columns: {len(df_reviews.dtypes)}')
print(f'Total Rows: {df_reviews.count():,}')
df_reviews.printSchema()
```

```
Total Columns: 9
Total Rows: 8,635,403
root
 |-- business_id: string (nullable = true)
 |-- cool: long (nullable = true)
 |-- date: string (nullable = true)
 |-- funny: long (nullable = true)
 |-- review_id: string (nullable = true)
 |-- stars: double (nullable = true)
 |-- text: string (nullable = true)
 |-- useful: long (nullable = true)
 |-- user_id: string (nullable = true)
```

Let's begin by listing the `business_id` and `stars` columns together for the user reviews data.

```
In [15]: # Show 2 columns of review dataframe
df_stars = df_reviews.select('business_id', 'stars')
df_stars.show(5)
```

```
+-----+-----+
|      business_id|stars|
+-----+-----+
|buF9druCkbuXLX526...| 4.0|
```

```
|RA4V8pr014UyUbDvI...| 4.0|
|_sS2LBIGNT5NQb6PD...| 5.0|
|0AzLzHf0JgL7R0whd...| 2.0|
|8zehGz9jnxPqXt0c7...| 4.0|
+-----+-----+
only showing top 5 rows
```

Now, let's aggregate along the `stars` column to get a resultant dataframe that displays *average stars* per business as accumulated by users who **took the time to submit a written review**.

```
In [16]: # Get average score of each business, only reviews with greater than 150 characters are counted
df_reviews.createOrReplaceTempView('Yelpdata_reviews')

sqldf = spark.sql(
'''
SELECT business_id, avg(stars) as Review_Score
FROM Yelpdata_reviews
WHERE length(text) >= 150
GROUP BY business_id
LIMIT 20
'''
)

sqldf.show(5)
```

```
+-----+-----+
|      business_id|      Review_Score|
+-----+-----+
|uEUweopM30lHcVxj0...|2.7777777777777777|
|wdBrDCbZopowEkIEX...| 4.466666666666667|
|bOnsvrz1VkbrZM1jV...|          3.8|
|R0IJhEI-zSJpYT1YN...| 3.533333333333333|
|XzXcpPCb8Y5huklEN...| 4.666666666666667|
+-----+-----+
only showing top 5 rows
```

Now the fun part - let's join our two dataframes (reviews and business data) by `business_id` .

```
In [17]: # join business and reviews dataframes

df_business.createOrReplaceTempView('Yelpdata_business')
```

```

sqldf1 = spark.sql(
'''
SELECT yb.business_id, avg(yr.stars) as Review_Score, first(yb.name) as name, first(yb.city) as city, first(yb.state) as
FROM Yelpdata_business yb
JOIN Yelpdata_reviews yr ON yb.business_id = yr.business_id
GROUP BY yb.business_id
LIMIT 20
'''
)

```

Let's see a few of these:

```
In [18]: sqldf1.show()
```

business_id	Review_Score	name	city	state
--0DF12EMHYI8XIgo...	4.333333333333333	Joe's Advanced Au...	Smyrna	GA
--hLwt1d-CymEs7KT...	3.833333333333333	Little Amsterdam ...	Milwaukie	OR
-0sIQ96u8XevGUXZ-...	3.779816513761468	Copacabana Cuba Cafe	Altamonte Springs	FL
-1WrFiy7X43HD40xF...	3.730769230769231	Cinemark	Columbus	OH
-1rvXk4zbX3I6ddMC...	4.087378640776699	ZenCha Tea Cafe	Bexley	OH
-2-RiF6h0SVVkjVwe...	2.857142857142857	Canton Plumbing &...	Canton	MA
-20FRpDsjwtFJP31W...	2.666666666666666	Denny's	PORTLAND	OR
-2joeHbqY9TayADes...	4.119760479041916	Lynwood Cafe	Randolph	MA
-3CNyvgW7ASXPmSZ...	2.75	Stanley Park Ghos...	Vancouver	BC
-3ZAPcamH9lPHvkb0...	3.846153846153846	The Westin Orland...	Orlando	FL
-497StkBbvRrxUs6R...	3.175675675675676	House of Vintage	Portland	OR
-4k7fspe6MGo06WpS...	4.0	Penrose Realty	Boston	MA
-5G8ujs_p2Ir1u6sN...	4.5	Elite Import Auto...	Austin	TX
-5Uqm5_grHV7LCdC1...	4.4	The Dance Inn	Lexington	MA
-5VyAi8GR34xmDagF...	3.700228832951945	Verde	Atlanta	GA
-75Dnp3pdz1Y1tBCE...	1.352941176470588	Sun Country Airlines	Portland	OR
-7jzk0NUJP2aFyERL...	2.5	PetSmart	Davenport	FL
-8N16tYqXP-ixDDzw...	4.0	The Coffee Bean &...	Austin	TX
-93SwBU3VREpx0hG0...	4.2	Td Bank Financial...	Vancouver	BC
-9AJAELrWVza80mgh...	4.5	North Shore Barte...	Salem	MA

Compute a new dataframe that calculates what we will call the *skew* (for lack of a better word) between the avg stars accumulated from

written reviews and the *actual* star rating of a business (ie: the average of stars given by reviewers who wrote an actual review **and** reviewers who just provided a star rating).

The formula you can use is something like:

$$(\text{row['avg(stars)']} - \text{row['stars']}) / \text{row['stars']}$$

If the **skew** is negative, we can interpret that to be: reviewers who left a written response were more dissatisfied than normal. If **skew** is positive, we can interpret that to be: reviewers who left a written response were more satisfied than normal.

In [19]:

```
# skew test

# (row['avg(stars)'] - row['stars']) / row['stars']

sqldf2 = spark.sql(
'''
SELECT t1.business_id, t1.Review_Score_with_text, t2.Total_Avg_score, (t1.Review_Score_with_text-t2.Total_Avg_score) as S
FROM
(
SELECT yb.business_id, avg(yr.stars) as Review_Score_with_text
FROM Yelpdata_business yb
JOIN Yelpdata_reviews yr ON yb.business_id = yr.business_id
WHERE length(yr.text) >= 150
GROUP BY yb.business_id ) t1

INNER JOIN

(SELECT yb.business_id, avg(yr.stars) as Total_Avg_Score
FROM Yelpdata_business yb
JOIN Yelpdata_reviews yr ON yb.business_id = yr.business_id
GROUP BY yb.business_id) t2

ON t1.business_id = t2.business_id

'''
)

sqldf2.show()
```

+-----+-----+-----+-----+

business_id	Review_Score_with_text	Total_Avg_score	Skew
--JuLhLvq3gyjNnXT...	5.0	5.0	0.0
--_nBudPOb1lNRgKf...	3.875	3.875	0.0
--kyOk0waSrCD1bSv...	3.7777777777777777	3.8666666666666667	-0.08888888888888902
--z9usx6Fin8P_f0v...	5.0	5.0	0.0
-0qeY1293steyCqYh...	3.375	3.375	0.0
-0wZIJnbYSstEGj3u...	1.8	1.8	0.0
-1Dcv3siosFTgDJhN...	4.076923076923077	4.2	-0.12307692307692353
-28M_3R-Iq3KjY3fa...	2.7	2.5454545454545454	0.15454545454545476
-2DAKRNW7SwJG4aKo...	3.689655172413793	3.774193548387097	-0.08453837597330383
-2ysHxktRcDom1m9A...	5.0	5.0	0.0
-3IqRwpj-iTz1nQWC...	3.1666666666666665	2.857142857142857	0.3095238095238093
-4DDszZHzSvY5unbt...	1.8113207547169812	1.8392857142857142	-0.02796495956873...
-5zKNFxuoPm8L9Q00...	4.344827586206897	4.34375	0.001077586206896...
-66HiDa701jBc2qK8...	3.6623376623376624	3.6744186046511627	-0.01208094231350021
-6FX2iidcEY50MOY_...	4.2444444444444445	4.2075471698113205	0.03689727463312398
-6QKHeA7pr2uLtqJd...	5.0	5.0	0.0
-6gIAeMUuR7XJVioB...	2.5762711864406778	2.693548387096774	-0.11727720065609626
-6uXqbL1TvTHLyfsS...	3.4	3.4	0.0
-777fa6pKF-dcIJLo...	3.3421052631578947	3.5777777777777778	-0.23567251461988326
-7C5aM1UytSeiHksX...	4.8	4.8	0.0

only showing top 20 rows

And finally, graph it!

In [20]:


```
# plot it
pandas_skew_df = sqldf2.toPandas()

plt.figure(figsize = (15,8))
plt.hist(pandas_skew_df['Skew'], density=True, bins = 100)

plt.margins(x=-.3, y=.25) # Values in (-0.5, 0.0) zooms in to center

plt.ylabel('Count')
plt.xlabel('Review Score Difference')
plt.show()

%matplotlib plt
```



So, do Yelp (written) Reviews skew negative? Does this analysis actually prove anything? Expound on implications / interpretations of this graph.

Yes, the scores skew negative. This means that reviewers who write long text reviews tend to be more dissatisfied with the business.

Extra Credit

Let's expand on the previous analysis by breaking out the skew into types of businesses. We will be looking at the 2 top businesses: restaurants and shopping.

```
In [21]: # Lets query the skew by categories  
# join with association table for categories  
  
df_assc.createOrReplaceTempView('Yelpdata_categories')
```

```

sqldf_ec = spark.sql(
...
SELECT

t1.business_id, yc.categories, t1.Review_Score_with_text, t2.Total_Avg_score,
(t1.Review_Score_with_text-t2.Total_Avg_score) as Skew

FROM

(SELECT yb.business_id, avg(yr.stars) as Review_Score_with_text
FROM Yelpdata_business yb
JOIN Yelpdata_reviews yr ON yb.business_id = yr.business_id
WHERE length(yr.text) >= 150
GROUP BY yb.business_id) t1

INNER JOIN

(SELECT yb.business_id, avg(yr.stars) as Total_Avg_Score
FROM Yelpdata_business yb
JOIN Yelpdata_reviews yr ON yb.business_id = yr.business_id
GROUP BY yb.business_id) t2 ON t1.business_id = t2.business_id

INNER JOIN

Yelpdata_categories yc ON t1.business_id = yc.business_id

WHERE yc.categories IN('Restaurants', 'Shopping', 'Food')

...
)

sqldf_ec.show()

```

business_id	categories	Review_Score_with_text	Total_Avg_score	Skew
--JuLhLvq3gyjNnXT...	Shopping	5.0	5.0	0.0
--_nBudPOb1lNRgKf...	Restaurants	3.875	3.875	0.0
--ky0k0waSrCDlbSv...	Food	3.7777777777777777	3.8666666666666667	-0.088888888888888902
--ky0k0waSrCDlbSv...	Restaurants	3.7777777777777777	3.8666666666666667	-0.088888888888888902
-1Dcv3siosFTgDJhN...	Food	4.076923076923077	4.2	-0.12307692307692353
-2ysHxktRcDom1m9A...	Shopping	5.0	5.0	0.0

-3IqRwpj-iTz1nQWC...	Restaurants	3.1666666666666665	2.857142857142857	0.3095238095238093
-5zKNFxuoPm8L9OQ0...	Shopping	4.344827586206897	4.34375	0.001077586206896...
-66HiDa701jBc2qK8...	Restaurants	3.6623376623376624	3.6744186046511627	-0.01208094231350021
-6FX2iidcEY5OMOY_...	Food	4.2444444444444445	4.2075471698113205	0.03689727463312398
-6FX2iidcEY5OMOY_...	Restaurants	4.2444444444444445	4.2075471698113205	0.03689727463312398
-6gIAeMUuR7XJVioB...	Shopping	2.5762711864406778	2.693548387096774	-0.11727720065609626
-777fa6pKF-dcIJLo...	Restaurants	3.3421052631578947	3.5777777777777778	-0.23567251461988326
-7Q-UJjq1LJw_2fzu...	Food	3.1666666666666665	3.1666666666666665	0.0
-7Q-UJjq1LJw_2fzu...	Restaurants	3.1666666666666665	3.1666666666666665	0.0
-7mViF5PiylD0tmH7...	Shopping	4.75	4.695652173913044	0.0543478260869561
-9ap9pStLtFBYOMLR...	Food	3.4591836734693877	3.532110091743119	-0.07292641827373147
-9ap9pStLtFBYOMLR...	Restaurants	3.4591836734693877	3.532110091743119	-0.07292641827373147
-DaGidyE_Vq15fQkB...	Shopping	5.0	5.0	0.0
-DqwYtCZMWJFk_9m0...	Shopping	2.6	2.3333333333333335	0.2666666666666666

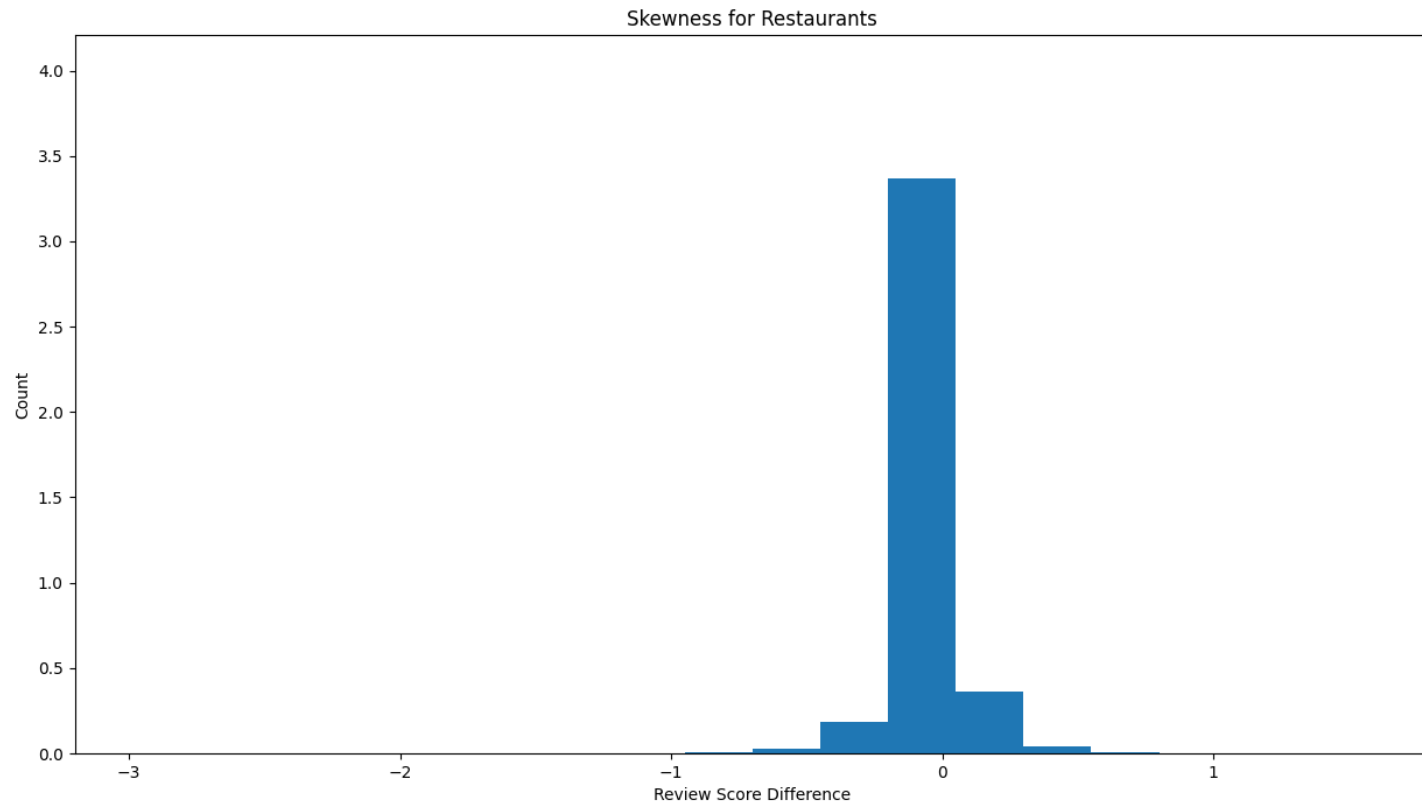
+-----+-----+-----+-----+-----+-----+
only showing top 20 rows

In [22]:

```
# visualize the skew by Restaurants
sqldf_ec
pandas_skew_categories = sqldf_ec.toPandas()
df_cate = pandas_skew_categories.loc[pandas_skew_categories['categories']=='Restaurants']

plt.figure(figsize = (15,8))
plt.hist(df_cate['Skew'], density=True, bins = 20)
plt.margins(x=0, y=.25)
plt.title("Skewness for Restaurants")
plt.ylabel('Count')
plt.xlabel('Review Score Difference')
plt.show()

%matplotlib plt
```

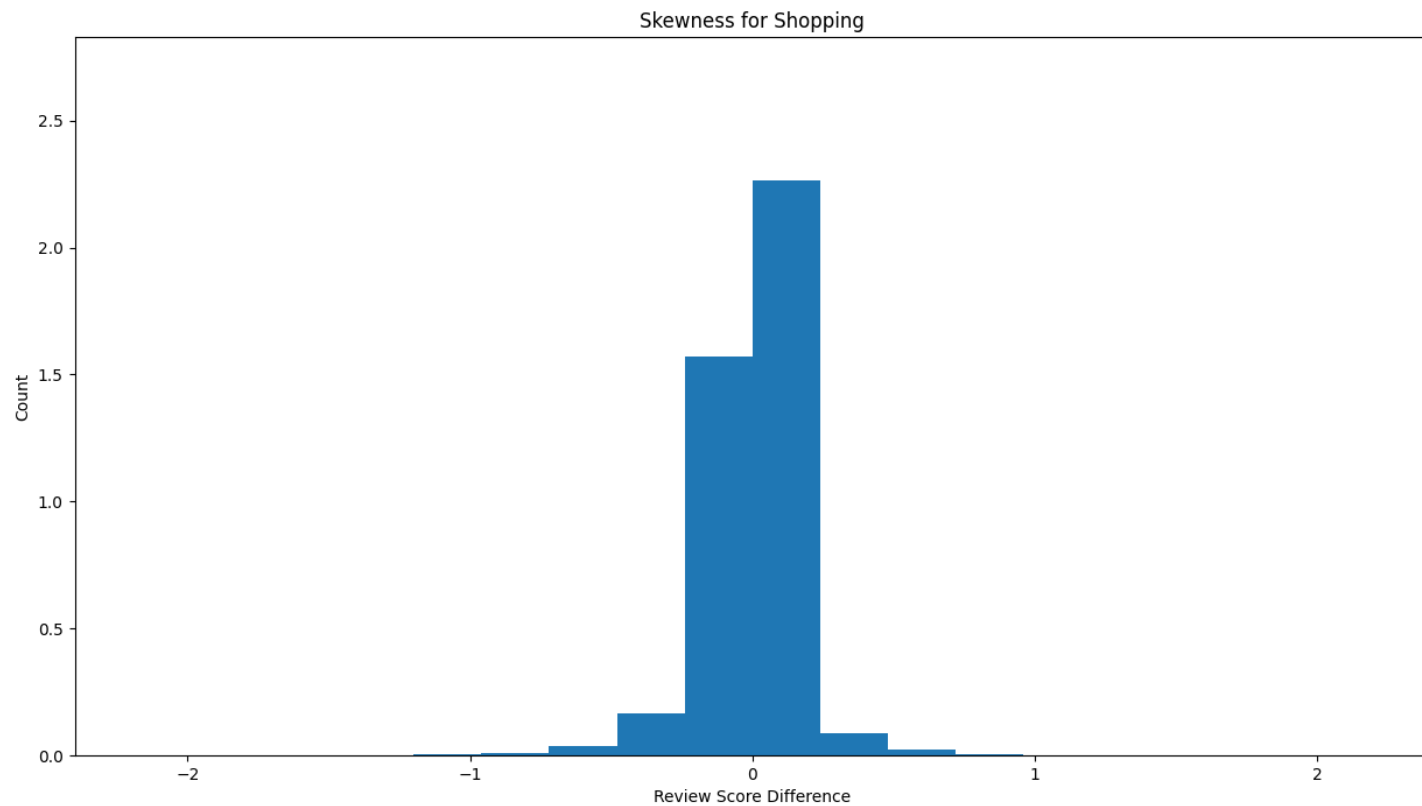


For restaurant reviews, review score with long text tend to be slightly lower than average total score

```
In [23]: # vizualize the skew by Shopping
sqlidf_ec
pandas_skew_categories = sqlidf_ec.toPandas()
df_cate = pandas_skew_categories.loc[pandas_skew_categories['categories']=='Shopping']

plt.figure(figsize = (15,8))
plt.hist(df_cate['Skew'], density=True, bins = 20)
plt.margins(x=0, y=.25)
plt.title("Skewness for Shopping")
plt.ylabel('Count')
plt.xlabel('Review Score Difference')
plt.show()
```

```
%matplotlib plt
```



No major noticable skewness for shopping reviews.

Should the Elite be Trusted? (Or, some other analysis of your choice)

For the final portion - you have a choice:

- Try and analyze some interesting dimension to this data. The **ONLY** requirement is that you must use the **Users** dataset and join on either the **business* or reviews**** dataset
- Or, you may try and answer the question posed: how accurate or close are the ratings of an "elite" user (check Users table schema) vs the actual business rating.

Feel free to use any and all methodologies at your disposal - only requirement is you must render one visualization in your analysis

Loading User Data

In [24]:

```
# Load the user data from s3
df_users = spark.read.json('s3://yelpdatasta9760/yelp_academic_dataset_user.json')

print(f'Total Columns: {len(df_users.dtypes)}')
print(f'Total Rows: {df_users.count():,}')
df_users.printSchema()
```

Total Columns: 22

Total Rows: 2,189,457

root

```
|-- average_stars: double (nullable = true)
|-- compliment_cool: long (nullable = true)
|-- compliment_cute: long (nullable = true)
|-- compliment_funny: long (nullable = true)
|-- compliment_hot: long (nullable = true)
|-- compliment_list: long (nullable = true)
|-- compliment_more: long (nullable = true)
|-- compliment_note: long (nullable = true)
|-- compliment_photos: long (nullable = true)
|-- compliment_plain: long (nullable = true)
|-- compliment_profile: long (nullable = true)
|-- compliment_writer: long (nullable = true)
|-- cool: long (nullable = true)
|-- elite: string (nullable = true)
|-- fans: long (nullable = true)
|-- friends: string (nullable = true)
|-- funny: long (nullable = true)
|-- name: string (nullable = true)
|-- review_count: long (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
|-- yelping_since: string (nullable = true)
```


In [25]:

```
# Show some columns to see what it looks like
df_show_users = df_users.select('user_id', 'name', 'review_count', 'elite', 'average_stars', 'yelping_since')
df_show_users.show(20)
```

user_id	name	review_count	elite	average_stars	yelping_since
q_QQ5kBBwLCcbL1s4...	Jane	1220	2006,2007,2008,20...	3.85	2005-03-14 20:26:35
dIIKEfOgo0KqUfGQv...	Gabi	2136	2007,2008,2009,20...	4.09	2007-08-10 19:01:51
D6ErcUnFALnQCN4b1...	Jason	119	2010,2011	3.76	2007-02-07 15:47:53
JnPIjvC0cmooNDfsa...	Kat	987	2009,2010,2011,20...	3.77	2009-02-09 16:14:29
37Hc8hr3cw0iHLoPz...	Christine	495	2009,2010,2011	3.72	2008-03-03 04:57:05
n-QwITZYrXlKQRiV3...	Natasha	229	2010,2011,2012,20...	3.59	2008-06-25 14:53:17
eCJoZqpV1fDKJGAsX...	Bridget	51		3.86	2009-07-22 16:47:01
cojec0wQJpsYDxnjt...	Steven	51	2010,2011	3.79	2010-07-04 17:18:40
1jXmzuIFKxTnEnR0p...	Clara	299	2010,2011,2012,20...	3.43	2010-10-01 17:29:36
-8Qo0IfvwxJ4sY20...	Antoinette	288	2012,2013,2014,20...	3.88	2007-08-04 20:21:09
EtofuImujQBS002xa...	Hollyanna	44	2009,2010	3.83	2008-11-03 17:31:30
cxS6dbjyPgPS1S890...	Joe	65		4.18	2005-07-22 11:07:20
MUzkXfPS9JaMgJ907...	Damon	750	2010,2012,2013,20...	4.16	2009-01-29 03:48:53
tjwblGkWN9m0vsGay...	Ben	632	2010,2011,2012,20...	3.66	2009-09-30 12:20:56
m-zIVssiXN4bnDFqM...	Bryan	2		4.5	2009-04-18 13:14:33
fxqvYXlml4400Bgls...	Cristina	363	2011,2012,2013,20...	3.65	2010-01-10 22:31:48
9edAbpniyhHFdpAvk...	Cara	356	2007,2008,2009,2010	3.73	2007-10-03 23:21:51
wURnB9fRNGAl1i3yB...	Jessie	77	2009,2010,2011,2012	3.75	2009-02-11 17:40:20
l4P65LXNBnJqI7oTX...	Lisa	202	2008,2009,2010,20...	3.93	2008-03-31 17:05:27
9RIXlhUb_xEVuc_o0...	silly	53		3.57	2009-02-28 17:37:36

only showing top 20 rows

In [26]:

```
# Build association table for easier analytics of user table
df_inter = df_users.withColumn('elite', psf.explode(psf.split("elite", ",")))
df_assc_user = df_inter.select('user_id', 'name', 'review_count', 'elite', 'average_stars', 'yelping_since')
df_assc_user.show(5)
```

user_id	name	review_count	elite	average_stars	yelping_since
q_QQ5kBBwLCcbL1s4...	Jane	1220	2006	3.85	2005-03-14 20:26:35
q_QQ5kBBwLCcbL1s4...	Jane	1220	2007	3.85	2005-03-14 20:26:35

q_QQ5kBBw1CcbL1s4...	Jane	1220	2008	3.85	2005-03-14 20:26:35
q_QQ5kBBw1CcbL1s4...	Jane	1220	2009	3.85	2005-03-14 20:26:35
q_QQ5kBBw1CcbL1s4...	Jane	1220	2010	3.85	2005-03-14 20:26:35

only showing top 5 rows

In [27]:

```
# Calculate Elite only review score and compare to Total Average score

df_assc_user.createOrReplaceTempView('Yelpdata_assc_users')

sqldf3 = spark.sql(
...
SELECT t1.business_id, t1.EliteOnly_Review_Score, t2.Total_Avg_Score, (t1.EliteOnly_Review_Score-t2.Total_Avg_Score) as E
FROM

(SELECT yr.business_id, avg(yr.stars) as EliteOnly_Review_Score
FROM Yelpdata_reviews yr
JOIN Yelpdata_assc_users yu ON (yu.user_id = yr.user_id AND YEAR(yr.date)= yu.elite)
GROUP BY yr.business_id
ORDER BY yr.business_id) t1

JOIN

(SELECT yb.business_id, avg(yr.stars) as Total_Avg_Score
FROM Yelpdata_business yb
JOIN Yelpdata_reviews yr ON yb.business_id = yr.business_id
GROUP BY yb.business_id
ORDER BY yb.business_id) t2

ON t1.business_id=t2.business_id

...
)

sqldf3.show()
```

business_id	EliteOnly_Review_Score	Total_Avg_Score	Elite_Skew
--164t1nclzzmca7e...	3.6451612903225805	3.877551020408163	-0.2323897300855826
--6COJIAjkQwSUZci...	3.9411764705882355	4.107843137254902	-0.16666666666666607
--DzGwfuJH12DjYz9...	3.0	2.2666666666666666	0.7333333333333334
--JKSSgnfoOjVDFGv...	5.0	4.04	0.96

--JuLhLvq3gyjNnXT...	5.0	5.0	0.0
--ToovR1Ob2e131Zi...	5.0	4.68	0.3200000000000003
--UNNdnHRhsyFUbDg...	4.4798850574712645	4.390995260663507	0.08888979680775755
--_nBudPOb1lNRgKf...	4.0	3.875	0.125
--bbZa1KPYSmW0X4o...	3.9	4.148936170212766	-0.2489361702127657
--g3lcoxHd8djBe1y...	4.0	4.230769230769231	-0.23076923076923084
--hLwt1d-CymEs7KT...	4.0	3.8333333333333335	0.16666666666666652
--hkbIWgBKBOZq4Vc...	4.474452554744525	4.301587301587301	0.1728652531572239
--iiZK8pw7kC2Tr2P...	2.8	3.1645569620253164	-0.3645569620253166
--jtHbM1FC1W0XIh9...	5.0	5.0	0.0
--kK4a5ggLXioTc4K...	3.736842105263158	3.6041666666666665	0.13267543859649145
--ky0k0waSrCD1bSv...	3.75	3.8666666666666667	-0.11666666666666667
--lXp7qjUuniGt3n-...	5.0	4.583333333333333	0.41666666666666696
--mvQZf74PNpgTKA4...	5.0	4.1666666666666667	0.8333333333333333
--oGYZhLMMvwfHmty...	3.0	2.933333333333333	0.06666666666666687
--vPNAA8d8UUruvu0...	3.6666666666666665	4.2631578947368425	-0.596491228070176

+-----+-----+-----+-----+
only showing top 20 rows

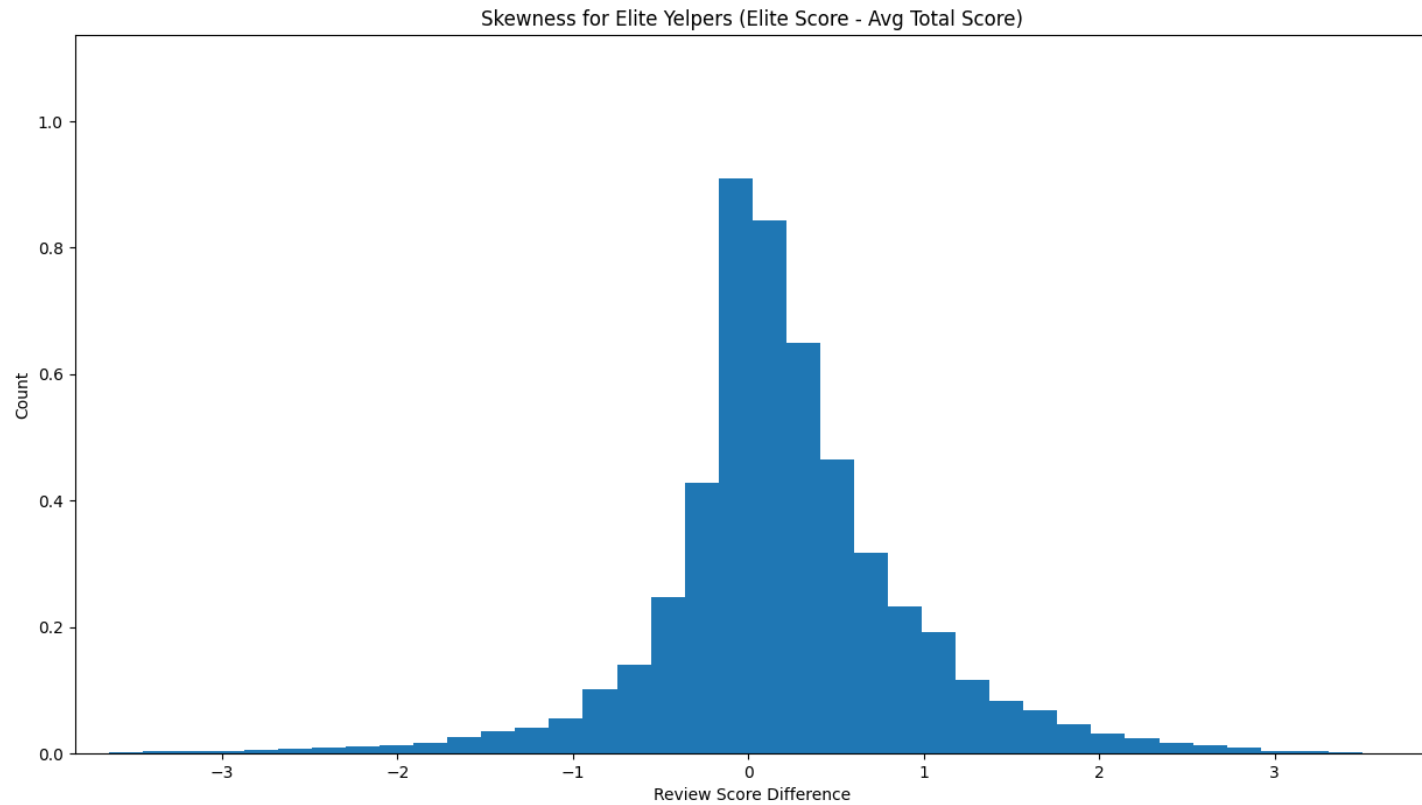
In [28]:

```
# plot it
elite_skew_df = sqldf3.toPandas()

plt.figure(figsize = (15,8))
plt.hist(elite_skew_df['Elite_Skew'], density=True, bins = 40)
plt.margins(x=0, y=.25)
plt.title("Skewness for Elite Yelpers (Elite Score - Avg Total Score)")
plt.ylabel('Count')
plt.xlabel('Review Score Difference')
plt.show()
```

In [29]:

```
%matplotlib plt
```



from the plot, elite users tend to give very slightly (almost not noticable) HIGHER score than the total average reviews.